

TreatmentFlow: Finding the Optimal Patient Triage System for Hospital Usage

Emese Elkind and Adam Neto

Abstract

Ontario emergency departments have an average wait time of 2.2 hours, with triage nurses responsible for prioritizing patients from levels 1–5 based on urgency—a role known to be highly stressful. To address this, we propose *TreatmentFlow*, an AI-based digital solution for optimizing the flow of patients through the emergency department. We introduce TreatmentFlow, an AI-driven system designed to improve emergency patient flow through automated triage, bed assignment, and diagnosis support. Our Deep Learning model predicts triage priority with 66% accuracy, showing promise but requiring improvement. For bed assignments, a Greedy Constraint Optimization approach efficiently minimizes delays, outperforming more complex constraint programming. A Bayesian Network approach achieves perfect accuracy in classifying pre-defined diagnoses, though real-time input limitations persist. Together, these modules aim to streamline ED operations, ease nurse workload, and assist physicians. While AI tools show clear potential in healthcare, challenges remain in accuracy and computational complexity.

I. Introduction

HEALTHCARE is constantly evolving. With emergency room wait times averaging 2.2 hours in Ontario [1], Artificial-Intelligence-driven triage systems could be the key to prioritizing patients more efficiently and reducing delays. Triage is the process hospitals use to determine the urgency of treatment for each patient [2]. This is determined by the patient's conditions and symptoms when they enter the hospital. Triage ensures that patients needing urgent care wait the least amount of time for treatment. In Ontario, patients are assigned a Canadian Triage and Acuity Scale (CTAS) priority number from 1-5 based on urgency of care (Table 1):

Table 1: Canadian Triage and Acuity Scale (CTAS) priority leveling [7].

CTAS Levels	Description
1	Resuscitation
2	Emergent
3	Urgent
4	Less Urgent
5	Non-Urgent

Triage nurses check vital signs and symptoms and gather patient history data to assess where they fit on this priority list. This is known to be one of the most stressful positions a nurse can have in the hospital [3]. Technology is increasingly being used in hospitals to assist with triage and relieve some of this stress on nurses with artificial intelligence (AI) and machine learning (ML) becoming novel topics for triage.

We propose *TreatmentFlow*, an AI-based solution for optimizing the flow of patients through the emergency room. We examine three problems in triage and apply AI techniques to solve them. Firstly, we aim to automate the assignment of patient priority using triage symptom information, relieving some triage decision-making pressure for nurses. The method presented to solve this is a Deep Learning approach. In this approach, we use a Deep Neural Network model to learn from data to be able to predict the CTAS priority level of incoming patients given symptom data.

Next, we seek to optimally assign patients to beds based on patient priority, further simplifying the decision-making tasks of the nurses. This is approached with two different methods: Mixed Integer Programming (MIP) Constraint Optimization, and Greedy patient assignment. In the MIP approach, we construct a model of variables and constraints to mimic a hospital and optimize (minimize) the total wait time of patients, weighted by priority level. This weighting implies that the model considers that the wait times of high-priority patients are more detrimental than the wait times of lower-priority patients. To test the value of this MIP solution, we compare its results to those of a greedy patient assignment algorithm. This comparison evaluates either solution's performance and efficiency at scale, including objective (sum of weighted wait times) minimization, and execution time.

Finally, our solution involves utilizing probabilistic inference and Bayesian Networks to produce each patient a probability table to display potential conditions given a set of triage symptoms. Doctors in the emergency department often ask for each other's opinions to ensure correct diagnosis. Double-checking diagnoses is important but can sometimes waste time and be ineffective. We propose a more efficient solution: to predict possible diagnoses and produce a doctors-only bedside document that includes the probability of predicted diagnoses, thus aiding in the efficient, correct diagnosis of patients. By providing the probabilities of patient conditions, doctors will be able to more efficiently deliberate between methods of treatment.

II. Methodology

Two datasets containing patient symptom information were used to train the Deep Learning and Bayesian Network components of *TreatmentFlow*: the *Hospital Triage and Patient History Data* from Yale New Haven Health System, covering March 2014 to July 2017, and *SymbiPredict Symptom-Disease Prediction Dataset (SDPD)*, sourced from Kaggle [4] and Mendeley [5], respectively.

The Yale dataset includes approximately 50,000 patient records with 972 variables, including patient information and symptoms. To manage file size and enable efficient analysis, the original 103.45 MB RData file was converted to CSV format and down sampled by selecting every 75th row, resulting in a reduced file size of 19.8 MB ensuring some level of random sampling while maintaining a manageable dataset size for analysis.




Figure 1: Distribution of CTAS values in the dataset.

The *SymbiPredict* dataset, comprising nearly 5,000 symptom–prognosis pairs in CSV format, required no preprocessing due to its smaller size (1.32 MB) and clean structure.

It is important to note that both datasets originate from U.S. healthcare systems, where Emergency Severity Index (ESI) is used for patient triage instead of the Canadian CTAS scale. Therefore, the python scripts reference ESI while this paper refers to CTAS for consistency with the Canadian healthcare context.

A. Constraint Satisfaction

The Constraint Satisfaction component of *TreatmentFlow* is a set of solutions for hospital bed assignment based on incoming patients and available resources. The two solutions are (1) a mixed integer programming (MIP) constraint optimization solution and (2) a greedy bed assignment solution. These solutions are compared in terms of efficiency and results on a scale.

Both solutions rely on a system of hospital objects formed in the file “HospitalClasses.py” these necessary objects include Bed, Patient, and HospitalRecords, each of which store and track information for instances of their category. For example, each Bed object stores their occupant (if any), each Patient stores statistics like CTAS priority level and arrival time, and HospitalRecords stores a database of patients. In our first iteration of *TreatmentFlow*’s constraint satisfaction component, we track the *inverse* of CTAS priorities, meaning we treat 5 as the highest priority and 1 as the lowest to simplify current calculations. We are also using a randomly generated set of patients, arrival times, and priorities, whereas we will chain all three components together for a given set of patients in our final product.

In the greedy solution to hospital bed assignment found in the “GreedyBedAssignment.py” file, we examine the current state of the hospital minute-by-minute over a 24-hour period of with the

assistance of custom HospitalRecords attributes such as “unserved” (the priority queue of unserved patients) and “serving” (the running list of patients occupying a bed). For each minute of the simulation, the scheduler assigns the next available highest-priority patient to any available bed and keeps them there for the duration of their service time. This assignment considers the patient’s CTAS priority level (inverse) and arrival time at the hospital. Given the number of patients (p) and number of beds (b), the hospital simulator method “run_hospital()” features a time complexity of $O(p * b)$, which is very efficient at scale. While this approach is algorithmically efficient, it is naïve in the fact that it does not guarantee an optimal solution but instead assigns patients based on currently available information.

Constraint optimization problems involve modeling a set of variables and assigning them values such that all constraints on variables are satisfied. They then use this model to optimize for a given objective function, which is either minimized or maximized. In the mixed integer programming (MIP) approach to hospital bed assignment found in the “MIPBedAssignment.py” file, we build a constraint satisfaction/optimization model using Google OR-Tools’ `cp_model` (Constraint Programming model) class. This library allows for the construction of constraint satisfaction models and provides a built-in solver that optimizes for a given objective function. Our model includes variables for each bed at each minute of the day, each of which can be set to any one of the patient identifications from the list of patients. The model builds several constraints concerning each patient: the arrival time constraint ensures that each patient cannot be assigned a bed before they arrive at the hospital; the contiguity constraint ensures that a patient does not shift between beds, and implicitly ensures that a patient can only be assigned to one bed at a time; and the tenure constraints ensure that a patient is not assigned to a bed for any shorter or longer than their service time, unless their service time extends past midnight (in which case they are assigned until midnight). It is noteworthy that the greedy assignment solution also satisfies all the same constraints. Finally, the MIP solution forms an objective function that minimizes the sum of the wait time of each patient (multiplied by inverse of CTAS priority level, which essentially translates to “patient cost”), applying a penalty of patient cost * total minutes per day (1440 min) to patients that are not assigned. When the solver is run, it finds solutions that satisfy all constraints, minimizing the objective function, and thus optimizing the bed assignment. We have also included user-entered parameters to set the number of beds, patients, and time limit of the MIP solver’s execution. Note that if an optimal solution is not found before the time limit runs out, the solver displays the most optimal solution found so far.


To test the value and validity of the MIP and Greedy solutions to bed assignment, we repeatedly compared the results of the same set of patients and beds. Across various scales, we performed this comparison with a new, random set of patients and beds at each denomination. The results are displayed in the next section of this report.

B. Deep Learning

To automate the assignment of patient priority using triage symptom information to relieve the pressure for nurses making wrong decisions a Deep Learning model assigns patient priority through this automated triage system.

Supervised machine learning has been chosen as the machine learning technique for this task because it uses labeled data to train algorithms to predict outcomes. The model is given labeled training to learn the relationship between the input and the outputs. The two main types of supervised machine learning are classification and regression [6]. Classification was chosen for this task to perform multi-class classification task, where the model predicts which of 5 possible priority levels (1-5) a patient should be assigned [7]. The output is the priority levels (1-5) which is discrete therefore the final layer of the algorithm will have 5 neurons. Common classification algorithms are linear classifiers, support vector machines (SVM), decision trees, k-nearest neighbor and random forest [6], Multi-layer Perceptron (MLP) classifier was chosen as a model for this dataset because it can handle mixed data types and is scalable for large datasets.

Hospital Triage and Patient History Data is used for the deep learning aspect of *TreatmentFlow*, with this data loaded into the file and preprocessed. The data is first observed using the `load_data` function, and Pandas [8] is used to observe the file format. A graph is made to observe the given CTAS scores, and a bell curve with level 3 being the most common priority level is shown. The data is preprocessed by dropping columns from the data set that are non-medical demographic data, as to not contribute to any data leakage. The data in the *Hospital Triage and Patient History Data* csv file presents itself with numerical and categorical features which is split using Numpy [9] to ensure that the model is trained on high-quality data. Data is split into 80% training data and 20% test data sets with the input represented by 'x' and are the columns of the file and the target output represented by 'y' which are the CTAS values. The MLP classification algorithm that is being used is a Deep Neural Network (DNN) and is implemented using Scikit Learn [11] with 3 hidden layers. To determine whether a neuron is activated or not the ReLU



activation function is used for the hidden layers. The ReLU activation function is the most common in deep learning and is a mathematical function to introduce non-linearity. If the input is positive the output is positive and if the input is negative the output is zero.

The model is evaluated using the testing data set (20% of the whole data set) which is saved in a confusion matrix to visualize the results.

C. Bayesian Networks

SymbiPredict encodes symptoms as binary values (1 for presence, 0 for absence) and a final column indicating the confirmed diagnosis. We implemented a supervised machine learning approach to classify diseases based on patient-reported binary symptom data. A Bernoulli Naive

Bayes classifier was chosen due to its suitability for binary feature data. The model was trained using the training set ($X_{\text{train}}, y_{\text{train}}$) and a default smoothing parameter $\alpha=1.0$, leveraging the BernoulliNB implementation from scikit-learn. Model performance was assessed on the test set using classification accuracy and detailed metrics (precision, recall, and F1-score) derived from the classification_report and confusion_matrix. In addition, 5-fold cross-validation was conducted on the full dataset to evaluate the model's generalizability. The mean and standard deviation of cross-validation scores were reported. To simulate clinical application, a prediction function was implemented that allows input of a dictionary of symptoms (present = 1, absent = 0). This function constructs a binary vector aligned with the training features and uses the trained model to predict the probability distribution across all known diseases. The top N most probable diseases are returned as a ranked list. To identify the most influential symptoms for a given disease, the log-probability values learned by the Bernoulli Naive Bayes model were exponentiated and sorted in descending order. This provided insight into which symptoms were most indicative of each condition, based on the conditional probability

$$P(\text{Diagnosis} \mid \text{Symptoms}) \propto P(\text{Diagnosis}) \times \prod_i P(\text{Symptom}_i \mid \text{Diagnosis})$$

Where $P(\text{Diagnosis} \mid \text{Symptoms})$ represents the posterior probability of a diagnosis given the observed symptoms, $P(\text{Diagnosis})$ represents the prior probability of the diagnosis, and $P(\text{Symptom}_i \mid \text{Diagnosis})$ represents the conditional probability of observing each symptom given the diagnosis.

The diagnostic report is a well formatted document containing the top few most likely diagnoses and state the observed symptoms that lead to those diagnoses.

III. Results

A. Constraint Satisfaction

Table 1: Ratio of objective values produced by the MIP approach and Greedy approach. For each cell, a ratio lower than zero means the MIP approach produced a better solution than the Greedy Approach within the 120 second time limit on the solver. A cell value of “None” means the MIP solver produced no solution. (x axis: Patients, y axis: Beds)

	5	10	20	50	100	250
5	1	1	1	2.6045	0.729	None
10	1	1	1	24344	None	None
15	1	1	1	inf	None	None
20	1	1	0	84078	None	None


Table 2: Ratio of execution times of MIP approach (with 120 second time limit on solver) and Greedy approach. (x axis: Patients, y axis: Beds)

	5	10	20	50	100	250
5	29015	33523	17991	58924	26433	18235
10	21085	59913	57062	100775	62713	24673
15	52333	67038	80641	111624	81349	34696
20	68041	88934	115251	119926	61192	45063

Table 3: Greedy approach’s patients’ average wait time in minutes, and average number of unserved patients over 100 attempts. (x axis: Patients, y axis: Beds)

	5	10	20	50	100	250	500
5	0, 0	0, 0	4, 0	51, 6	39, 66	33, 221	24, 471
10	0, 0	0, 0	0, 0	0, 0	40, 11	24, 188	33, 441
25	0, 0	0, 0	0, 0	0, 0	0, 0	30, 21	24, 345
50	0, 0	0, 0	0, 0	0, 0	0, 0	0, 0	24, 34
100	0, 0	0, 0	0, 0	0, 0	0, 0	0, 0	0, 0

B. Deep Learning



C. Bayesian Networks

MIP solving is notably faster when the number of patients is close to or lower than the number of beds, as the increased solution space allows for more feasible optimal assignments. This is because there are more possible solutions that allow an optimal assignment of patients to beds.

This chance for optimality comes from the fact that there is far more space at each individual minute, meaning that each patient is more likely to find an empty bed at their arrival time, and thus the MIP solver finds an optimal solution much faster. However, as the MIP solver more easily finds an optimal solution, the Greedy approach also has fewer calculations to perform, which stabilizes the ratio of either approach's execution time in table 3.

When comparing each solution's performance in terms of their objective values, the two solutions often perform similarly, with some randomization involved. This randomization of relative performance depends on the layout of patient arrival times in each randomized list of patients. As discussed above, the MIP solver finds optimal solutions much more often when the number of patients is near the number of beds, or equal, or lower. However, the greedy solution provides near optimal solutions in these scenarios as well in a fraction of the time. As shown in table 1, the Greedy approach produced an equal or superior objective value in all but two observations. This shows that as problem size increases, the computational inefficiency of the MIP solution begins to outweigh the potential improvements it has in bed assignment optimization. This makes it impractical for large-scale settings, and thus impractical for most hospitals.

To conclude the comparison, the MIP model can sometimes find better solutions than the Greedy method, but the drastically larger computational complexity of the model makes it infeasible at scale. Since the differences between solutions found are marginal when limiting the solver's execution time, the Greedy bed assignment method, despite being a naïve solution, remains the superior bed assignment method. Using the Greedy bed assignment tool, we can reduce the logistical delays of patient wait times to negligible values, assuming a realistic number of patients compared to capacity (Table 3). The functional limit for the ratio of patients to beds that allows this tool to minimize wait times is when the ratio is less than ten. Since this is a very extreme potential ratio of patients to hospital resources, this tool remains very useful, even at scale. Thus, this application of constraint satisfaction creates a significant benefit for the operation of emergency departments.

B. Deep Learning

The confusion matrix (Figure 2) is a table that is used to define the performance of a classification algorithm by comparing its predicted labels to the actual labels [10]. The y-axis of the matrix represents the true class, which are the actual labels, and the x-axis represents the predicted class, which are the model's output. The diagonal values from (1,1) to (5,5) represent the correct predictions, where the model's predicted priority matches the actual priority. The higher these values and the darker the coloring of the box, the better the model's accuracy. The values off the diagonal represent misclassifications, where the model predicted the wrong priority level. If a value is above the diagonal, it means the model underestimated priority. If a value is below the diagonal, it means the model overestimated priority. The model accuracy is 66% when trained on

hospital data from our data set. The model seems to have the most difficulty distinguishing between priority 1 and 2, as well as 2 and 3, based on the number of misclassifications between them. There are very few cases in priority 0 and 4, which might indicate class imbalance. A good DNN model accuracy is generally considered to be above 70% according to industry standards, with anything between 70% and 90% being considered a realistic and valuable performance range. The *TreatmentFlow* model is not up to industry standards yet, therefore, future work includes finding the optimal hidden layer nodes to maximize model accuracy and thus patient priority levels.

C. Bayesian Networks

The model accuracy on test set is 100% when trained on data from our data set. The confusion matrix (Figure 3) shows that only correct diagnoses are made with matches only on the diagonal values of the matrix which represent the correct predictions, where the model's predicted priority matches the actual priority. Even with a Naive Bayes classifier accuracy of perfect, there are observed limitations in real time diagnosis based on input symptoms. The observed limitations arise from missed symptoms inputted by the user.

A clinical diagnostic assessment was performed using our Bayesian Networks diagnostic tool on a simulated patient case (Patient ID: 1234). The system analyzed reported symptoms such as chest pain, fatigue, vomiting, and breathlessness, and generated a differential diagnosis based on symptom importance and presence. The algorithm identified Heart Attack as the primary diagnosis with 100% probability and a “Very High” confidence level. This determination was strongly supported by key symptoms like chest pain (importance: 0.9490), vomiting (0.9184), and breathlessness (0.8878), all of which were present in the patient. Alternative conditions are ruled out due to absence of other critical symptoms. Although some overlapping symptoms can be present across differential diagnoses, the algorithm accurately weighted their relative importance and cross-referenced them with the complete symptom profile. It is evident that the patient - doctor communication must be strong for the doctor to input the correct symptoms in order to have a strong diagnosis.

The output structure includes a ranked list of differential diagnoses with calculated probabilities, a table of symptom importance and presence, rationale for excluding other diagnoses. This structured and interpretable output not only supports clinical decision-making but also illustrates the transparency of our model. Such capabilities are critical for building trust for the doctors and physicians in our AI-assisted diagnosis tools.

V. Conclusion

AI-driven triage systems like *TreatmentFlow* have the potential to revolutionize emergency room efficiency by addressing key challenges in patient prioritization, bed assignment, and

diagnostic support. By leveraging Deep Learning for automated triage, Constraint Optimization for bed allocation, and Bayesian Networks for diagnostic assistance, *TreatmentFlow* aims to reduce emergency wait times, alleviate stress on healthcare workers, and improve patient outcomes. As hospitals continue to integrate AI technologies, solutions like this can play a critical role in optimizing resource management and ensuring timely, effective care for all patients.

The three parts of this AI project are tightly integrated: the output of one module becomes the input for the next. Specifically the CTAS priority levels predicted in the Deep Learning section passed into the Constraint Optimization section to optimally assign these patients to beds, and once the beds are assigned, it will be passed into the Bayesian Network section to predict diagnoses for a doctor-aimed bedside document. Aligning all three of these components should allow for a seamless assistive tool for healthcare workers to address patients' movement and treatment throughout the hospital.

In future iterations, *TreatmentFlow* can be validated through simulation environments and prospective clinical trials. Success for this project would be defined by its ability to outperform traditional triage in both speed and accuracy, reduce average emergency department wait times through patient prioritization and routing, and improve diagnostic precision. The system aims to ease the workload of clinical staff by offering automated, interpretable recommendations that support rather than replace human judgment.

Further work will also focus on enhancing the Deep Learning triage model, refining the optimization algorithms for dynamic constraints, and improving interpretability of diagnostic outputs to support clinical adoption and trust.

References

- [1] *System performance*. Emergency Department Time Spent by Patients in Ontario – Health Quality Ontario (HQO). (n.d.). <https://www.hqontario.ca/system-performance/time-spent-in-emergency-departments>
- [2] Peta, D., Day, A., Lugari, W., Gorman, V., Ahayalimudin, N., Pajo, V. (2023, November 2). *Triage: A global perspective*. Journal of Emergency Nursing. <https://www.sciencedirect.com/science/article/abs/pii/S0099176723002143#:~:text=Triage%20is%20a%20process%20by,timely%20access%20to%20emergency%20care.>
- [3] Emergency room triage: How does it work? - the daily scan. (n.d.). <https://thedailyscan.providencehealthcare.org/2018/11/emergency-room-triage-how-does-it-work/>
- [4] Mafaufau, M. (2019, June 3). *Hospital triage and patient history data*. Kaggle. <https://www.kaggle.com/datasets/maalona/hospital-triage-and-patient-history-data>

- [5] Tucker, J. (2024, April 2). *Symbipredict*. Mendeley Data.
<https://data.mendeley.com/datasets/dv5z3v2xyd/1>
- [6] Belcic, I., & Stryker, C. (2024, December 26). *What is supervised learning?*. IBM.
<https://www.ibm.com/think/topics/supervised-learning#:~:text=its%20training%20data.,Types%20of%20supervised%20learning,next%20layer%20in%20the%20network>.
- [7] *Emscimprovement*. Emergency Nurses Association. (2002).
https://media.emscimprovement.center/documents/Emergency_Severity_Index_Handbook.pdf
- [8] *Pandas*. pandas. (n.d.). <https://pandas.pydata.org/>
- [9] *NumPy*. (n.d.). <https://numpy.org/>
- [10] *Confusion matrix*. Confusion Matrix - an overview | ScienceDirect Topics. (n.d.).
<https://www.sciencedirect.com/topics/engineering/confusion-matrix#:~:text=A%20confusion%20matrix%20is%20a,malignant%20tissue%20is%20considered%20cancerous>.
- [11] *Learn*. scikit. (n.d.). <https://scikit-learn.org/stable/>